

Linh Huynh

San Jose, CA | vanlinhtpnt@gmail.com | (408) 690-5492 | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

EXPERIENCE

AI Intern | Helix International — Grand Rapids, MI

Jun 2025 – Sep 2025

- Benchmarked and deployed multiple open-source LLMs (Llama, Phi, GPT-oss, Qwen3) locally via vLLM on GPU desktops for offline inference; built 70+ benchmark scenarios across SQL-generation and natural-language tasks to select production-ready models.
- Ran load/performance tests with Locust simulating concurrent request traffic to measure latency, throughput, and failure rates across models.
- Built a React chat interface integrated with a FastAPI + LangGraph backend using RAG to convert SQL results into plain-English answers.
- Raised answer accuracy above 90% via prompt engineering, temperature tuning, and post-processing/context-narrowing to reduce hallucination.

CS Tutoring | UC Davis — Davis, CA

Jan 2025 – Jun 2025

- Mentored **10+ students** in Java, Python, and core computer science topics including data structures and algorithms.
- Guided students through debugging, implementation strategies, and software engineering best practices while improving problem-solving skills.

TECHNICAL SKILLS

Languages: Python, Java, C++, SQL, HTML/CSS, JavaScript

Backend: FastAPI, Node.js, Express, REST APIs

Frontend: React, Typescript, Tailwind CSS, HTML/CSS, D3.js, Recharts

AI / ML: LLM Applications, LangGraph, Retrieval-Augmented Generation (RAG), Prompt Engineering, OpenAI API, Gemini API, PyTorch

Databases: PostgreSQL, Supabase

Clouds & Tools: Git, GitHub, Docker, Vercel, Render, Ollama, vLLM

Work Authorization: F-1 OPT (STEM OPT Eligible)

PROJECTS

Job Search Automation Agent — Python, Anthropic Claude API, Pydantic

Aug 2026

- Architected a multi-stage AI agent pipeline (job discovery → requirement extraction → resume matching → tailored drafting) with a human-in-the-loop review gate and a code-enforced no-fabrication guarantee that verifies every AI-generated claim against source data.
- Diagnosed and fixed a quadratic API-cost blowup from a growing tool-call conversation, redesigning it into isolated single-purpose calls and cutting measured cost per job ~95% (\$1.88 → \$0.09).
- Built a deterministic verification layer that rejects any AI-claimed match unless the cited text is a verbatim substring of the source — eliminating a class of hallucination at the code level rather than via prompting.

JourZy — Full-Stack Travel Planner

Feb 2026 – Present

- Built and deployed a conversational web app generating personalized itineraries; integrated **Google Maps, Distance Matrix, and live weather/transit APIs** for real-time recommendations.
- Engineered structured LLM output (deterministic JSON parsing, inline rationales) using **Gemini and local Ollama/Qwen** models.
- Iterated with **5+ early users** (friends/family); deployed on **Vercel using Claude Code** in the workflow.

AI Operations Copilot — FastAPI, React, PostgreSQL, Recharts

Jun 2025 – Present

- Developed a **full-stack observability** platform for analyzing application logs and monitoring system health.
- Designed scalable backend pipelines for processing streaming log data and supporting real-time analytics.
- Implemented semantic search and AI-assisted debugging features to accelerate production incident investigation.
- Built interactive dashboards for visualizing error trends, timelines, and operational metrics.

EDUCATION

University of California, Davis | B.S. Computer Science

Graduated in June 2026

Relevant Coursework: Data Structures & Algorithms, Algorithm Design & Analysis, Operating Systems, Machine Learning, Computer Architecture, Computer Vision